

mmdetection源码剖析(1)--NMS

ManWingloeng 2020-07-04 原文

mmdetection源码剖析(1)--NMS

熟悉目标检测的应该都清楚**NMS**是什么算法，但是如果我们要与C++和cuda结合直接写成Pytorch的操作你们清楚怎么写吗？最近在看**mmdetection**的源码，发现其实原来写C++和cuda的扩展也不难，下面给大家讲一下。

C++的扩展是允许用户来创建自定义PyTorch框架外的操作（operators）的，即从PyTorch后端分离。此方法与实现本地PyTorch操作的方式不同。C++扩展旨在为您节省大量与将操作与PyTorch后端集成在一起相关的样板，同时为基于PyTorch的项目提供高度的灵活性。

官方给出了一个LLTM的例子，大家也可以看一下。

NMS算法

先复习一下NMS的算法：

- （1）将所有框的得分排序，选中最高分及其对应的框
- （2）遍历其余的框，如果和当前最高分框的重叠面积(IOU)大于一定阈值，我们就将框删除。
- （3）从未处理的框中继续选一个得分最高的，重复上述过程。

这里我给出一份纯numpy的实现：

```
1. def nms(bounding_boxes, Nt):
2.     if len(bounding_boxes) == 0:
3.         return [], []
4.     bboxes = np.array(bounding_boxes)
5.     x1 = bboxes[:, 0]
6.     y1 = bboxes[:, 1]
7.     x2 = bboxes[:, 2]
8.     y2 = bboxes[:, 3]
9.     scores = bboxes[:, 4]
10.    areas = (x2 - x1 + 1) * (y2 - y1 + 1)
11.    order = np.argsort(scores)
12.    picked_boxes = []
13.    while order.size > 0:
14.        index = order[-1]
15.        picked_boxes.append(bounding_boxes[index])
16.        x1l = np.maximum(x1[index], x1[order[:-1]])
17.        y1l = np.maximum(y1[index], y1[order[:-1]])
18.        x2l = np.minimum(x2[index], x2[order[:-1]])
19.        y2l = np.minimum(y2[index], y2[order[:-1]])
20.        w = np.maximum(0.0, x2l - x1l + 1)
21.        h = np.maximum(0.0, y2l - y1l + 1)
22.        intersection = w * h
```

你是有价值的 你很重要

当您需要与人交谈时 今天与我们交谈吧

辅导事工

编写Pytorch C++扩展的步骤

需要编写下面5个文件:

nms_kernel.cu

nms_cuda.cpp

nms_ext.cpp

setup.py

nms_wrapper.py

(1) nms_kernel.cu 主要使用ATen和THC库编写nms_cuda_forward的函数, 使用C++编写, 涉及一些lazyInitCUDA, THCudaFree, THCCeilDiv 的操作, 算法跟我们前面写的numpy差不多。

(2) nms_cuda.cpp 是调用了nms_kernel.cu文件的nms_cuda_forward封装了一下变成nms_cuda函数。

(3) nms_ext.cpp 进一步封装nms_cuda函数为nms, 并且通过PYBIND11_MODULE绑定成python可调用的函数。

```
1. PYBIND11_MODULE(TORCH_EXTENSION_NAME, m) {
2.     m.def("nms", &nms, "non-maximum suppression");
3. }
```

通过上面那样就相当于告诉python函数名定义为nms了。

(4) setup.py 就是编译一遍nms_ext, 至此你可以通过nms_ext.nms调用cpp extension作为pytorch的操作了

```
1. make_cuda_ext(
2.     name='nms_ext',
3.     module='mmdet.ops.nms',
4.     sources=['src/nms_ext.cpp', 'src/cpu/nms_cpu.cpp'],
5.     sources_cuda=[
6.         'src/cuda/nms_cuda.cpp', 'src/cuda/nms_kernel.cu'
7.     ]),
```

(5) nms_wrapper.py 再次封装 nms_ext.nms, 方便使用, 使用实例:

```
1. from . import nms_ext
2. inds = nms_ext.nms(dets_th, iou_thr)
```

稍微完整的代码如下, 但是我也删减了一些, 只剩下nms相关的代码, 想要看完整代码可以点击下面的文件名。

nms_kernel.cu (这个估计有部分是Facebook写的)

```
1. // Copyright (c) Facebook, Inc. and its affiliates. All Rights Reserved.
2. #include <ATen/ATen.h>
3. #include <ATen/cuda/CUDAContext.h>
4. #include <ATen/DeviceGuard.h>
5. #include <THC/THC.h>
6. #include <THC/THCDeviceUtils.cuh>
7. #include <vector>
8. #include <iostream>
9.
10. int const threadsPerBlock = sizeof(unsigned long long) * 8;
11.
12. __device__ inline float devIoU(float const * const a, float const * const b) {
13.     float left = max(a[0], b[0]), right = min(a[2], b[2]);
14.     float top = max(a[1], b[1]), bottom = min(a[3], b[3]);
15.     float width = max(right - left, 0.f), height = max(bottom - top, 0.f);
16.     float interS = width * height;
```

你是有价值的 你很重要

当您需要与人交谈时 今天与我们交谈吧

辅导事工

```

23.         const float *dev_boxes, unsigned long long *dev_mask) {
24.     const int row_start = blockIdx.y;
25.     const int col_start = blockIdx.x;
26.     // if (row_start > col_start) return;
27.     const int row_size =
28.         min(n_boxes - row_start * threadsPerBlock, threadsPerBlock);
29.     const int col_size =
30.         min(n_boxes - col_start * threadsPerBlock, threadsPerBlock);
31.
32.     __shared__ float block_boxes[threadsPerBlock * 5];
33.     if (threadIdx.x < col_size) {
34.         block_boxes[threadIdx.x * 5 + 0] =
35.             dev_boxes[(threadsPerBlock * col_start + threadIdx.x) * 5 + 0];
36.         block_boxes[threadIdx.x * 5 + 1] =
37.             dev_boxes[(threadsPerBlock * col_start + threadIdx.x) * 5 + 1];
38.         block_boxes[threadIdx.x * 5 + 2] =
39.             dev_boxes[(threadsPerBlock * col_start + threadIdx.x) * 5 + 2];
40.         block_boxes[threadIdx.x * 5 + 3] =
41.             dev_boxes[(threadsPerBlock * col_start + threadIdx.x) * 5 + 3];
42.         block_boxes[threadIdx.x * 5 + 4] =
43.             dev_boxes[(threadsPerBlock * col_start + threadIdx.x) * 5 + 4];
44.     }
45.     __syncthreads();
46.
47.     if (threadIdx.x < row_size) {
48.         const int cur_box_idx = threadsPerBlock * row_start + threadIdx.x;
49.         const float *cur_box = dev_boxes + cur_box_idx * 5;
50.         int i = 0;
51.         unsigned long long t = 0;
52.         int start = 0;
53.         if (row_start == col_start) {
54.             start = threadIdx.x + 1;
55.         }
56.         for (i = start; i < col_size; i++) {
57.             if (devIoU(cur_box, block_boxes + i * 5) > nms_overlap_thresh) {
58.                 t |= 1ULL << i;
59.             }
60.         }
61.         const int col_blocks = THCCeilDiv(n_boxes, threadsPerBlock);
62.         dev_mask[cur_box_idx * col_blocks + col_start] = t;
63.     }
64. }
65. // boxes is a N x 5 tensor
66. at::Tensor nms_cuda_forward(const at::Tensor boxes, float nms_overlap_thresh) {
67.     // Ensure CUDA uses the input tensor device.
68.     at::DeviceGuard guard(boxes.device());
69.     using scalar_t = float;
70.     AT_ASSERTM(boxes.device().is_cuda(), "boxes must be a CUDA tensor");
71.     auto scores = boxes.select(1, 4);
72.     auto order_t = std::get<1>(scores.sort(0, /* descending=*/true));
73.     auto boxes_sorted = boxes.index_select(0, order_t);
74.
75.     int boxes_num = boxes.size(0);
76.
77.     const int col_blocks = THCCeilDiv(boxes_num, threadsPerBlock);
78.
79.     scalar_t* boxes_dev = boxes_sorted.data_ptr<scalar_t>();
80.
81.     THCState *state = at::globalContext().lazyInitCUDA(); // TODO replace with getTHCState
82.
83.     unsigned long long* mask_dev = NULL;
84.     //THCudaCheck(THCudaMalloc(state, (void**) &mask_dev,
85.     //                boxes_num * col_blocks * sizeof(unsigned long long)));
86.
87.     mask_dev = (unsigned long long*) THCudaMalloc(state, boxes_num * col_blocks * sizeof(unsigned
long long));
88.
89.     dim3 blocks(THCCeilDiv(boxes_num, threadsPerBlock),
90.                 THCCeilDiv(boxes_num, threadsPerBlock));
91.     dim3 threads(threadsPerBlock);
92.     nms_kernel<<<blocks, threads, 0, at::cuda::getCurrentCUDAStream()>>>(boxes_num,
93.                                     nms_overlap_thresh,

```

你是有价值的 你很重要

当您需要与人交谈时 今天与我们交谈吧

辅导事工

```

100.         mask_dev,
101.         sizeof(unsigned long long) * boxes_num * col_blocks,
102.         cudaMemcpyDeviceToHost,
103.         at::cuda::getCurrentCUDASTream()
104.     ));
105.
106.     std::vector<unsigned long long> remv(col_blocks);
107.     memset(&remv[0], 0, sizeof(unsigned long long) * col_blocks);
108.
109.     at::Tensor keep = at::empty({boxes_num}, boxes.options().dtype(at::kLong).device(at::kCPU));
110.     int64_t* keep_out = keep.data_ptr<int64_t>();
111.
112.     int num_to_keep = 0;
113.     for (int i = 0; i < boxes_num; i++) {
114.         int nblock = i / threadsPerBlock;
115.         int inblock = i % threadsPerBlock;
116.
117.         if (!(remv[nblock] & (1ULL << inblock))) {
118.             keep_out[num_to_keep++] = i;
119.             unsigned long long *p = &mask_host[0] + i * col_blocks;
120.             for (int j = nblock; j < col_blocks; j++) {
121.                 remv[j] |= p[j];
122.             }
123.         }
124.     }
125.     THCudaFree(state, mask_dev);
126.     // TODO improve this part
127.     return order_t.index({
128.         keep.narrow(/*dim=*/0, /*start=*/0, /*length=*/num_to_keep).to(
129.             order_t.device(), keep.scalar_type())});
130. }

```

nms_cuda.cpp

```

1.  #include <torch/extension.h>
2.  #define CHECK_CUDA(x) TORCH_CHECK(x.device().is_cuda(), #x, " must be a CUDATensor ")
3.  at::Tensor nms_cuda_forward(const at::Tensor boxes, float nms_overlap_thresh);
4.  at::Tensor nms_cuda(const at::Tensor& dets, const float threshold) {
5.      CHECK_CUDA(dets);
6.      if (dets.numel() == 0)
7.          return at::empty({0}, dets.options().dtype(at::kLong).device(at::kCPU));
8.      return nms_cuda_forward(dets, threshold);
9.  }

```

nms_ext.cpp

```

1.  #include <torch/extension.h>
2.  #ifdef WITH_CUDA
3.  at::Tensor nms_cuda(const at::Tensor& dets, const float threshold);
4.  #endif
5.  at::Tensor nms(const at::Tensor& dets, const float threshold) {
6.      if (dets.device().is_cuda()) {
7.          #ifdef WITH_CUDA
8.              return nms_cuda(dets, threshold);
9.          #else
10.             AT_ERROR("nms is not compiled with GPU support");
11.          #endif
12.      }
13.      # return nms_cpu(dets, threshold);
14.  }
15.  PYBIND11_MODULE(TORCH_EXTENSION_NAME, m) {
16.      m.def("nms", &nms, "non-maximum suppression");
17.  }

```

setup.py

```

1.  def make_cuda_ext(name, module, sources, sources_cuda=[]):

```

你是有价值的 你很重要

当您需要与人交谈时 今天与我们交谈吧

辅导事工

```

8.     extension = CUDAEExtension
9.     extra_compile_args['nvcc'] = [
10.         '-D__CUDA_NO_HALF_OPERATORS__',
11.         '-D__CUDA_NO_HALF_CONVERSIONS__',
12.         '-D__CUDA_NO_HALF2_OPERATORS__',
13.     ]
14.     sources += sources_cuda
15. else:
16.     print(f'Compiling {name} without CUDA')
17.     extension = CppExtension
18.     # raise EnvironmentError('CUDA is required to compile MMDetection!')
19.
20.     return extension(
21.         name=f'{module}.{name}',
22.         sources=[os.path.join(*module.split('.'), p) for p in sources],
23.         define_macros=define_macros,
24.         extra_compile_args=extra_compile_args)
25.
26. if __name__ == '__main__':
27.     write_version_py()
28.     setup(
29.         name='mmdet',
30.         version=get_version(),
31.         description='Open MMLab Detection Toolbox and Benchmark',
32.         long_description=readme(),
33.         author='OpenMMLab',
34.         author_email='chenkaidev@gmail.com',
35.         keywords='computer vision, object detection',
36.         url='https://github.com/open-mmlab/mmdetection',
37.         packages=find_packages(exclude=('configs', 'tools', 'demo')),
38.         package_data={'mmdet.ops': ['*/*.so']},
39.         classifiers=[
40.             'Development Status :: 4 - Beta',
41.             'License :: OSI Approved :: Apache Software License',
42.             'Operating System :: OS Independent',
43.             'Programming Language :: Python :: 3',
44.             'Programming Language :: Python :: 3.5',
45.             'Programming Language :: Python :: 3.6',
46.             'Programming Language :: Python :: 3.7',
47.         ],
48.         license='Apache License 2.0',
49.         setup_requires=parse_requirements('requirements/build.txt'),
50.         tests_require=parse_requirements('requirements/tests.txt'),
51.         install_requires=parse_requirements('requirements/runtime.txt'),
52.         extras_require={
53.             'all': parse_requirements('requirements.txt'),
54.             'tests': parse_requirements('requirements/tests.txt'),
55.             'build': parse_requirements('requirements/build.txt'),
56.             'optional': parse_requirements('requirements/optional.txt'),
57.         },
58.         ext_modules=[
59.             make_cuda_ext(
60.                 name='compiling_info',
61.                 module='mmdet.ops.utils',
62.                 sources=['src/compiling_info.cpp']),
63.             make_cuda_ext(
64.                 name='nms_ext',
65.                 module='mmdet.ops.nms',
66.                 sources=['src/nms_ext.cpp', 'src/cpu/nms_cpu.cpp'],
67.                 sources_cuda=[
68.                     'src/cuda/nms_cuda.cpp', 'src/cuda/nms_kernel.cu'
69.                 ]),
70.         ],
71.         cmdclass={'build_ext': BuildExtension},
72.         zip_safe=False)

```

nms_wrapper.py

```

1. from . import nms_ext
2. def nms(dets, iou_thr, device_id=None):

```

你是有价值的 你很重要

当您需要与人交谈时 今天与我们交谈吧

辅导事工

```
9.         device = 'cpu' if device_id is None else f'cuda:{device_id}'
10.         dets_th = torch.from_numpy(dets).to(device)
11.     else:
12.         raise TypeError('dets must be either a Tensor or numpy array, '
13.                           f'but got {type(dets)}')
14.     # execute cpu or cuda nms
15.     if dets_th.shape[0] == 0:
16.         inds = dets_th.new_zeros(0, dtype=torch.long)
17.     else:
18.         if dets_th.is_cuda:
19.             inds = nms_ext.nms(dets_th, iou_thr)
20.         else:
21.             inds = nms_ext.nms(dets_th, iou_thr)
22.
23.     if is_numpy:
24.         inds = inds.cpu().numpy()
25.     return dets[inds, :], inds
```

总结

想要编写一个c++和cuda的扩展给Pytorch使用，其实主要就4步：

使用ATEN和THC编写前向代码cu文件A

封装成一个cpp文件B

再把B封装一遍并且使用PYBIND11_MODULE绑定函数名function

通过make_cuda_ext（这个是mmdetection自定义的函数）把组件setup一遍

通过绑定的函数名function就可以在Pytorch中调用了。

mmdetection源码剖析(1)--NMS的更多相关文章

1. jQuery之Deferred源码剖析

一.前言 大约在夏季,我们谈过ES6的Promise(详见here),其实在ES6前jQuery早就有了Promise,也就是我们所知道的Deferred对象,宗旨当然也和ES6的Promise一样, ...

2. Nodejs事件引擎libuv源码剖析之：高效线程池(threadpool)的实现

声明:本文为原创博文,转载请注明出处. Nodejs编程是全异步的,这就意味着我们不必每次都阻塞等待该次操作的结果,而事件完成(就绪)时会主动回调通知我们.在网络编程中,一般都是基于Reactor线程 ...

3. Apache Spark源码剖析

Apache Spark源码剖析(全面系统介绍Spark源码,提供分析源码的实用技巧和合理的阅读顺序,充分了解Spark的设计思想和运行机理) 许鹏 著 ISBN 978-7-121-25420- ...

4. 基于mybatis-generator-core 1.3.5项目的修订版以及源码剖析

项目简单说明 mybatis-generator,是根据数据库表.字段反向生成实体类等代码文件.我在国庆时候,没事剖析了mybatis-generator-core源码,写了相当详细的中文注释,可以去 ...

5. STL"源码"剖析-重点知识总结

STL是C++重要的组件之一,大学时看过<STL源码剖析>这本书,这几天复习了一下,总结出以下LZ认为比较重要的知识点,内容有点略多 :) 1.STL概述 STL提供六大组件,彼此可以组合 ...

6. SpringMVC源码剖析（四）- DispatcherServlet请求转发的实现

SpringMVC完成初始化流程之后,就进入Servlet标准生命周期的第二个阶段,即"service"阶段.在"service"阶段中,每一次Http请求到来,容器都会启动一个请求线程,通过serv ...

7. 自己实现多线程的socket，socketserver源码剖析

1.IO多路复用 三种多路复用的机制:select.poll.epoll 用的多的两个:select和epoll 简单的说就是:1.select和poll所有平台都支持,epoll只有linux支持2 ...